

DSE3002 - Project 3

May 1, 2023

```
[ ]: #DSE3002 - Project 3 - Patrick Norcross
```

```
[ ]: import pandas as pd
import numpy as np
from sklearn.metrics import accuracy_score, confusion_matrix, ConfusionMatrixDisplay

def AccConMax(y_test, y_pred):
    print(accuracy_score(y_test, y_pred))
    cm = confusion_matrix(y_test, y_pred)
    ConfusionMatrixDisplay(cm, display_labels=['Low', 'Normal', 'High']).plot()
```

```
[ ]: #1 - Importing data from
data18 = pd.read_csv('ATUS18_3.csv')
data19 = pd.read_csv('ATUS19_3.csv')
data21 = pd.read_csv('ATUS21_3.csv')
```

```
[ ]: #Data set up (all precrossing done in project 2 code)
X_18 = data18.iloc[:, :len(data18.columns)-1].values
y_18 = data18.iloc[:, len(data18.columns)-1:].values

y_18 = y_18.reshape(len(y_18),)
```

```
[ ]: X_19 = data19.iloc[:, :len(data19.columns)-1].values
y_19 = data19.iloc[:, len(data19.columns)-1:].values

y_19 = y_19.reshape(len(y_19),)
```

```
[ ]: X_21 = data21.iloc[:, :len(data21.columns)-1].values
y_21 = data21.iloc[:, len(data21.columns)-1:].values

y_21 = y_21.reshape(len(y_21),)
```

PART 1

Question 2

```
[ ]: from sklearn.ensemble import BaggingClassifier, AdaBoostClassifier,
      RandomForestClassifier
from tabulate import tabulate

[ ]: #Bagging

#Lower max features were not performing as well usually in the .50 range for
#the score.
#Below are some examples of parameters I tested.
scoreBag = 0
list = []
maxs = [.75,.25,.25,.5,.75]
maxf = [.75,.75,1.,1.,1.]
for s, f in zip(maxs,maxf):
    for est in [50,100,200]:
        bagx = BaggingClassifier(n_estimators=est, max_samples= s,
        max_features= f, oob_score=True, random_state=1)
        bagx.fit(X_18, y_18)
        list.append([s,f,est, bagx.oob_score_])
        if bagx.oob_score_ > scoreBag:
            scoreBag = bagx.oob_score_
            bag = bagx

print(tabulate(list, headers=['sample prop','feature prop','estimators',
        'scoreBag']))
print('\nBest Bagging Model:\n', bag)
```

sample prop	feature prop	estimators	scoreBag
0.75	0.75	50	0.613579
0.75	0.75	100	0.616723
0.75	0.75	200	0.625105
0.25	0.75	50	0.609388
0.25	0.75	100	0.617351
0.25	0.75	200	0.613998
0.25	1	50	0.627619
0.25	1	100	0.637678
0.25	1	200	0.638935
0.5	1	50	0.633697
0.5	1	100	0.647318
0.5	1	200	0.650042
0.75	1	50	0.633068
0.75	1	100	0.639355
0.75	1	200	0.639774

Best Bagging Model:

BaggingClassifier(max_samples=0.5, n_estimators=200, oob_score=True,

```
random_state=1)
```

```
[ ]: #Random Forest
import warnings
warnings.filterwarnings('ignore')

scoreRF = 0
list = []
for maxf in ['auto', None]:
    for est in [10,50,100,200]:
        rfx = RandomForestClassifier(n_estimators= est, max_features= maxf,
        oob_score=True, random_state=1)
        rfx.fit(X_18, y_18)
        list.append([maxf,est, rfx.oob_score_])
        if rfx.oob_score_ > scoreRF:
            scoreRF = rfx.oob_score_
            rf = rfx

print(tabulate(list, headers=['max feat setting', 'num of feat', 'score']))
print('\nBest RF Model:\n', rf)
```

max feat setting	num of feat	score
auto	10	0.495599
auto	50	0.595138
auto	100	0.604359
auto	200	0.615884
	10	0.539816
	50	0.629925
	100	0.637888
	200	0.640821

Best RF Model:

```
RandomForestClassifier(max_features=None, n_estimators=200, oob_score=True,
                        random_state=1)
```

```
[ ]: #AdaBoost
scoreBoost = 0
list = []
for est in [50,100,200]:
    for lr in [0.01,0.1,0.5,1,2]:
        boostx = AdaBoostClassifier(n_estimators=est, learning_rate=lr,
        random_state=1)
        boostx.fit(X_18, y_18)
        list.append([est, lr, boostx.score(X_18,y_18)])
        if boostx.score(X_18, y_18) > scoreBoost:
            scoreBoost = boostx.score(X_18, y_18)
            boost = boostx
```

```
print(tabulate(list, headers=['estimator', 'learning rate', 'score']))
print('\nBest Adaboost Model:\n', boost)
```

estimator	learning rate	score
50	0.01	0.503982
50	0.1	0.549874
50	0.5	0.606873
50	1	0.651718
50	2	0.333822
100	0.01	0.503982
100	0.1	0.567687
100	0.5	0.644174
100	1	0.682733
100	2	0.333822
200	0.01	0.503982
200	0.1	0.593043
200	0.5	0.667854
200	1	0.723177
200	2	0.333822

Best Adaboost Model:

```
AdaBoostClassifier(learning_rate=1, n_estimators=200, random_state=1)
```

Question 3

```
[ ]: print('RF:      ', scoreRF)
      print('Bagging: ', scoreBag)
      print('Boost:   ', scoreBoost)

#Based on the out-of-bag score/score, I would expect the Boosting classifier to
→ predict the best since it is higher than best bagging and boosting.
#This may be because it is using a different scoring method.
```

RF: 0.6408214585079631

Bagging: 0.6500419111483655

Boost: 0.7231768650461022

Question 4

```
[ ]: #2019 CONFUSION MATRICES

y_19_pred_bag = bag.predict(X_19)
y_19_pred_rf = rf.predict(X_19)
y_19_pred_ada = boost.predict(X_19)

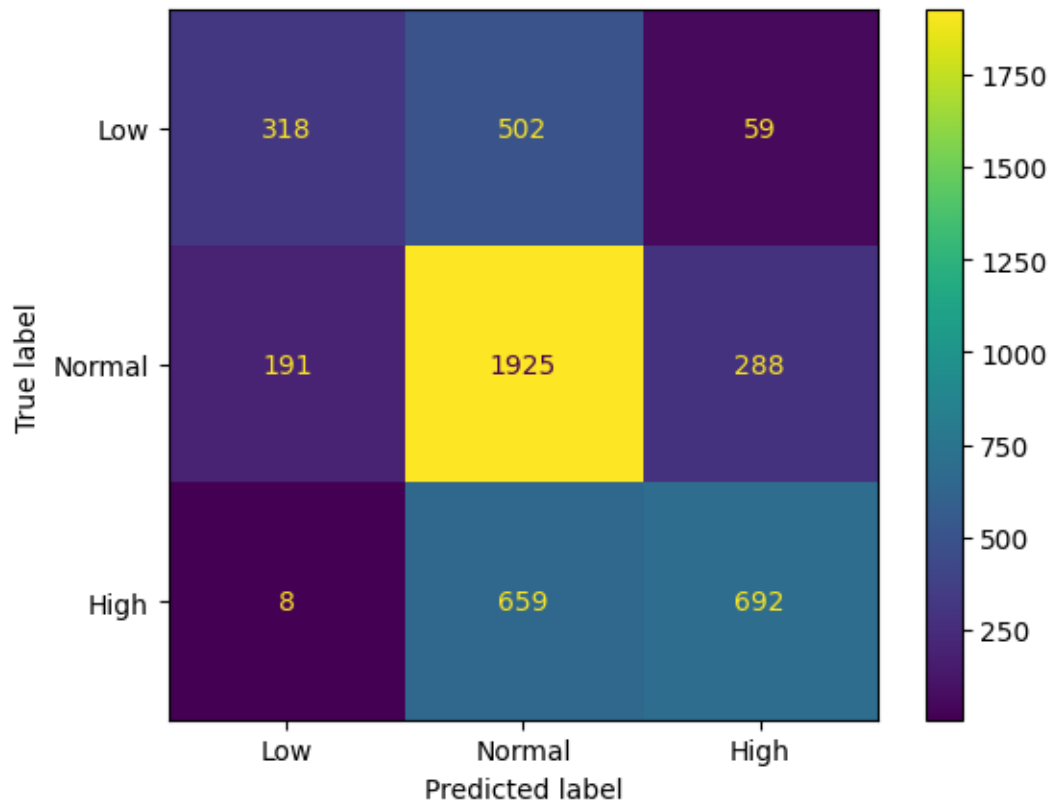
print('Random Forest')
AccConMax(y_19, y_19_pred_rf)
```

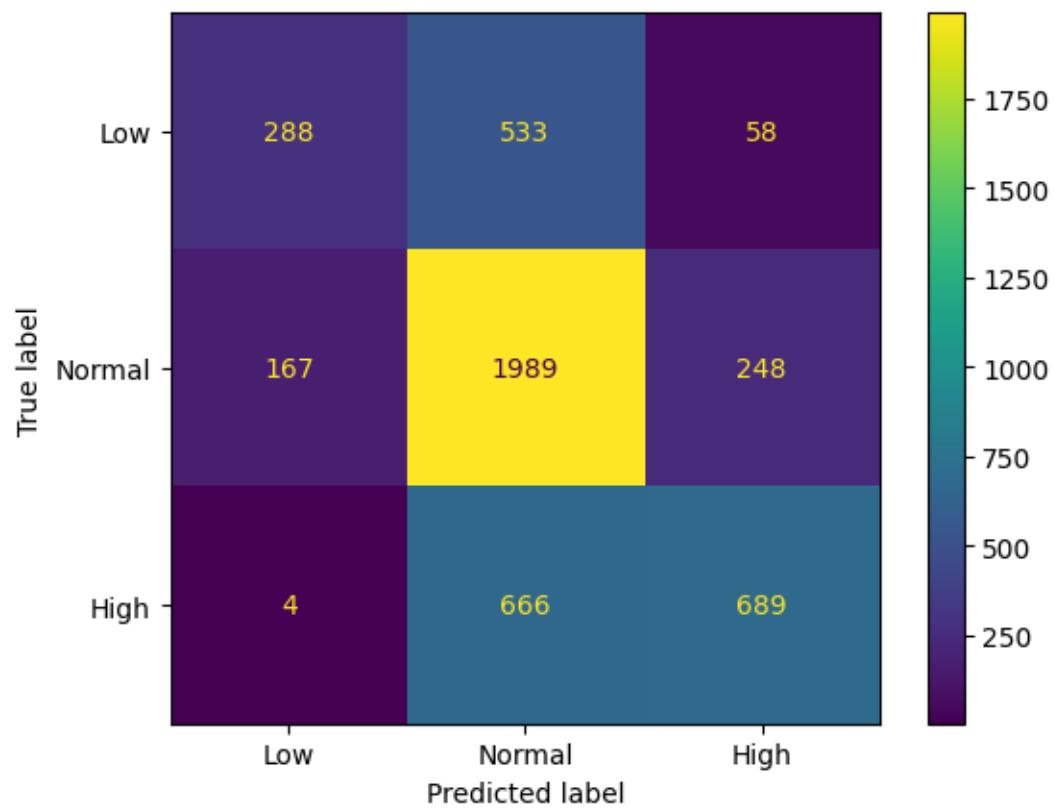
```


print('Bagging')
AccConMax(y_19, y_19_pred_bag)
print('ADA Boosting')
AccConMax(y_19, y_19_pred_ada)

```

Random Forest
 0.6322705730288669
 Bagging
 0.6389487289961223
 ADA Boosting
 0.6650150797070228







```
[ ]: #The ADABOOST did the best for 2019, which did match my hypothesis from step 3.
      #ADABOOST was biased toward Normal Sleep whereas the other 2 were focused more
      ↳ toward higher sleep levels.
```


```
[ ]: #2021 CONFUSION MATRICES
```


```
y_21_pred_bag = bag.predict(X_21)
y_21_pred_rf = rf.predict(X_21)
y_21_pred_ada = boost.predict(X_21)


print('Random Forest')
AccConMax(y_21, y_21_pred_rf)
print('Bagging')
AccConMax(y_21, y_21_pred_bag)
print('ADA Boosting')
AccConMax(y_21, y_21_pred_ada)
```

```
Random Forest
0.640464199693453
Bagging
0.6533829647470988
```

ADA Boosting
0.6822859645281366







```
[ ]: #The same things held in the 2021 data as the 2019 data.
      #The ADABOOST did the best for 2021, which did match my hypothesis from step 3.
      #ADABOOST was biased toward Normal Sleep whereas the other 2 were focused more
      ↳ toward higher sleep levels.
```

PART 2

Question 5

```
[ ]: from sklearn.tree import DecisionTreeClassifier
      from sklearn.neighbors import KNeighborsClassifier
      from sklearn.svm import SVC
      from sklearn.naive_bayes import GaussianNB
      from sklearn.ensemble import VotingClassifier
      from sklearn.model_selection import GridSearchCV
```

```
[ ]: tuple(range(1,11))
```

```
[ ]: (1, 2, 3, 4, 5, 6, 7, 8, 9, 10)
```

```
[ ]: #Decision Tree GridSearch
      parameters = {'max_depth': (1,2,3,4,5,6,7,8,9,10)}
```

```

GS = GridSearchCV(DecisionTreeClassifier(), parameters)
GS.fit(X_18, y_18)
dt = GS.best_estimator_
print(dt)
print(dt.score(X_18, y_18))

#kNN GridSearch
parameters = {'n_neighbors': (9, 13, 17, 19, 21, 23, 25, 27)} #not always the
↳best result
GS = GridSearchCV(KNeighborsClassifier(), parameters)
GS.fit(X_18, y_18)
knn = GS.best_estimator_
print(knn)
print(knn.score(X_18, y_18))

#SVC GridSearch
parameters = {'kernel': ('linear', 'rbf'), 'C': [10]} #I checked 1-10 but to
↳save on run time I set it to 10
svc = SVC(probability=True)
GS = GridSearchCV(svc, parameters)
GS.fit(X_18, y_18)
svm = GS.best_estimator_
print(svm)
print(svm.score(X_18, y_18))

#Naive Bayes (No GridSearch)
nb = GaussianNB()
nb.fit(X_18, y_18)
print(nb)
print(nb.score(X_18, y_18))

#SVM is the highest score at .87 while naivebayes is the lowest with .51

```

```

DecisionTreeClassifier(max_depth=8)
0.6823134953897737
KNeighborsClassifier(n_neighbors=21)
0.5752305113160101
SVC(C=10, kernel='linear', probability=True)
0.875314333612741
GaussianNB()
0.5132020117351216

```

```

[ ]: Vclf = VotingClassifier(estimators=[('dt',dt), ('knn',knn), ('svm',svm),
↳('nb',nb)],
                             voting='soft', weights=[1,1,1,1])

Vclf.fit(X_18, y_18)

```

```
Vclf.score(X_18, y_18)
```


```
[ ]: 0.7518860016764459
```

```
[ ]: Vclf = VotingClassifier(estimators=[('dt',dt),('knn',knn),('svm',svm),('nb',  
    ↪nb)],  
    voting='soft',weights=[3,1,6,0.5])  
  
Vclf.fit(X_18, y_18)  
Vclf.score(X_18, y_18)  
#This score is unsurprisingly higher than the one with equal weights since I  
    ↪prioritized the better performing models.
```

```
[ ]: 0.8883067896060352
```

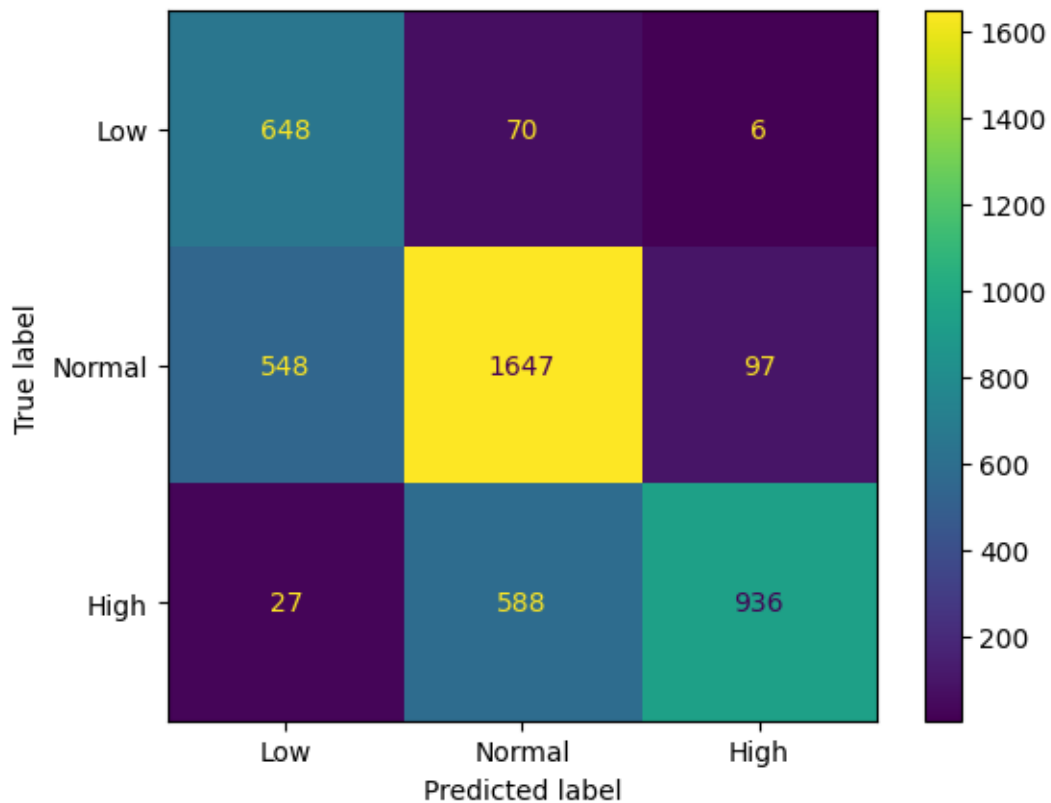
```
[ ]: #2019 CONFUSION MATRIX  
y_19_pred_vote = Vclf.predict(X_19)  
AccConMax(y_19, y_19_pred_vote)
```

```
0.6949590693666523
```



```
[ ]: #2021 CONFUSION MATRIX
y_21_pred_vote = Vclfw.predict(X_21)
AccConMax(y_21, y_21_pred_vote)
```

0.7074666082767681



```
[ ]: #Best 2019 - Voting Classifier
#The voting classifier does lower the accuracy of normal sleep but puts more
priority into both low and high sleep making it better for overall accuracy.
```

```
[ ]: #Best 2021 - Voting Classifier
#The voting classifier again lowers the accuracy of normal sleep but allows low
sleep and high sleep to have higher accuracy as well.
```